

Notes on AIMove exercises

Sylvain Calinon, Idiap Research Institute

1 Forward kinematics (FK) for a planar robot manipulator

The forward kinematics of a planar robot manipulator is defined as

$$\mathbf{f} = \begin{bmatrix} \mathbf{l}^\top \cos(\mathbf{L}\mathbf{x}) \\ \mathbf{l}^\top \sin(\mathbf{L}\mathbf{x}) \\ \mathbf{1}^\top \mathbf{x} \end{bmatrix} = \begin{bmatrix} l_1 \cos(x_1) + l_2 \cos(x_1 + x_2) + l_3 \cos(x_1 + x_2 + x_3) + \dots \\ l_1 \sin(x_1) + l_2 \sin(x_1 + x_2) + l_3 \sin(x_1 + x_2 + x_3) + \dots \\ x_1 + x_2 + x_3 + \dots \end{bmatrix},$$

with \mathbf{x} the state of the robot (joint angles), \mathbf{f} the position of the robot end-effector, \mathbf{l} a vector of robot links lengths, \mathbf{L} a lower triangular matrix with unit elements, and $\mathbf{1}$ a vector of unit elements.

The position and orientation of all articulations can similarly be computed with the forward kinematics function

$$\tilde{\mathbf{f}} = \begin{bmatrix} \mathbf{L} \text{diag}(\mathbf{l}) \cos(\mathbf{L}\mathbf{x}), & \mathbf{L} \text{diag}(\mathbf{l}) \sin(\mathbf{L}\mathbf{x}), & \mathbf{L}\mathbf{x} \end{bmatrix}^\top = \begin{bmatrix} \tilde{f}_{1,1} & \tilde{f}_{1,2} & \tilde{f}_{1,3} & \dots \\ \tilde{f}_{2,1} & \tilde{f}_{2,2} & \tilde{f}_{2,3} & \dots \\ \tilde{f}_{3,1} & \tilde{f}_{3,2} & \tilde{f}_{3,3} & \dots \end{bmatrix},$$

with

$$\begin{aligned} \tilde{f}_{1,1} &= l_1 \cos(x_1), \\ \tilde{f}_{2,1} &= l_1 \sin(x_1), \\ \tilde{f}_{3,1} &= x_1, \\ \tilde{f}_{1,2} &= l_1 \cos(x_1) + l_2 \cos(x_1 + x_2), \\ \tilde{f}_{2,2} &= l_1 \sin(x_1) + l_2 \sin(x_1 + x_2), \\ \tilde{f}_{3,2} &= x_1 + x_2, \\ \tilde{f}_{1,3} &= l_1 \cos(x_1) + l_2 \cos(x_1 + x_2) + l_3 \cos(x_1 + x_2 + x_3), \\ \tilde{f}_{2,3} &= l_1 \sin(x_1) + l_2 \sin(x_1 + x_2) + l_3 \sin(x_1 + x_2 + x_3), \\ \tilde{f}_{3,3} &= x_1 + x_2 + x_3, \\ &\vdots \end{aligned}$$

In Python, this can be coded for the end-effector position part as

```
1 D = 3 #State space dimension (joint angles)
2 x = np.ones(D) * np.pi / D #Robot pose
3 l = np.array([2, 2, 1]) #Links lengths
4 L = np.tril(np.ones([D,D])) #Transformation matrix
5 f = np.array([L @ np.diag(l) @ np.cos(L @ x), L @ np.
               diag(l) @ np.sin(L @ x)]) #Forward kinematics
```

2 Inverse kinematics (IK) for a planar robot manipulator

The Jacobian corresponding to the end-effector forward kinematics function can be computed as (with a simplification for the orientation part by ignoring the manifold aspect)

$$\mathbf{J} = \begin{bmatrix} -\sin(\mathbf{L}\mathbf{x})^\top \text{diag}(\mathbf{l})\mathbf{L} \\ \cos(\mathbf{L}\mathbf{x})^\top \text{diag}(\mathbf{l})\mathbf{L} \\ \mathbf{1}^\top \end{bmatrix} = \begin{bmatrix} J_{1,1} & J_{1,2} & J_{1,3} & \dots \\ J_{2,1} & J_{2,2} & J_{2,3} & \dots \\ J_{3,1} & J_{3,2} & J_{3,3} & \dots \end{bmatrix},$$

with

$$\begin{aligned} J_{1,1} &= -l_1 \sin(x_1) - l_2 \sin(x_1 + x_2) - l_3 \sin(x_1 + x_2 + x_3) - \dots, \\ J_{2,1} &= l_1 \cos(x_1) + l_2 \cos(x_1 + x_2) + l_3 \cos(x_1 + x_2 + x_3) + \dots, \\ J_{3,1} &= 1, \\ J_{1,2} &= -l_2 \sin(x_1 + x_2) - l_3 \sin(x_1 + x_2 + x_3) - \dots, \\ J_{2,2} &= l_2 \cos(x_1 + x_2) + l_3 \cos(x_1 + x_2 + x_3) + \dots, \\ J_{3,2} &= 1, \\ J_{1,3} &= -l_3 \sin(x_1 + x_2 + x_3) - \dots, \\ J_{2,3} &= l_3 \cos(x_1 + x_2 + x_3) + \dots, \\ J_{3,3} &= 1, \\ &\vdots \end{aligned}$$

In Python, this can be coded for the end-effector position part as

```
1 J = np.array([-np.sin(L @ x).T @ np.diag(l) @ L, np.cos(
               L @ x).T @ np.diag(l) @ L]) #Jacobian (for end-
               effector)
```

Numerical estimation of the Jacobian

Section 2 presents an analytical solution for the Jacobian. A numerical solution can alternatively be estimated by computing

$$J_{i,j} = \frac{\partial f_i(\mathbf{x})}{\partial x_j} \approx \frac{f_i(\mathbf{x}^{(j)}) - f_i(\mathbf{x})}{\Delta} \quad \forall i, \forall j,$$

with $\mathbf{x}^{(j)}$ a vector of the same size as \mathbf{x} with elements

$$x_k^{(j)} = \begin{cases} x_k + \Delta, & \text{if } k = j, \\ x_k, & \text{otherwise,} \end{cases}$$

where Δ is a small value for the approximation of the derivatives.

3 Linear quadratic tracking (LQT)

The LQT problem is formulated as the minimization of the cost

$$\begin{aligned} c &= (\boldsymbol{\mu}_T - \mathbf{x}_T)^\top \mathbf{Q}_T (\boldsymbol{\mu}_T - \mathbf{x}_T) \\ &\quad + \sum_{t=1}^{T-1} \left((\boldsymbol{\mu}_t - \mathbf{x}_t)^\top \mathbf{Q}_t (\boldsymbol{\mu}_t - \mathbf{x}_t) + \mathbf{u}_t^\top \mathbf{R}_t \mathbf{u}_t \right) \\ &= (\boldsymbol{\mu} - \mathbf{x})^\top \mathbf{Q} (\boldsymbol{\mu} - \mathbf{x}) + \mathbf{u}^\top \mathbf{R} \mathbf{u}, \end{aligned} \quad (1)$$

with $\mathbf{x} = [\mathbf{x}_1^\top, \mathbf{x}_2^\top, \dots, \mathbf{x}_T^\top]^\top$ the evolution of the state variable and $\mathbf{u} = [\mathbf{u}_1^\top, \mathbf{u}_2^\top, \dots, \mathbf{u}_{T-1}^\top]^\top$ the evolution of the control variable. $\boldsymbol{\mu} = [\boldsymbol{\mu}_1^\top, \boldsymbol{\mu}_2^\top, \dots, \boldsymbol{\mu}_T^\top]^\top$ represents the evolution of the tracking target. $\mathbf{Q} = \text{blockdiag}(\mathbf{Q}_1, \mathbf{Q}_2, \dots, \mathbf{Q}_T)$ represents the evolution of the required tracking precision, and $\mathbf{R} = \text{blockdiag}(\mathbf{R}_1, \mathbf{R}_2, \dots, \mathbf{R}_{T-1})$ represents the evolution of the cost on the control inputs.

The evolution of the system is linear, described by $\mathbf{x}_{t+1} = \mathbf{A}_t \mathbf{x}_t + \mathbf{B}_t \mathbf{u}_t$, yielding $\mathbf{x} = \mathbf{S}_x \mathbf{x}_1 + \mathbf{S}_u \mathbf{u}$ at trajectory level, see Appendix B for details.

The solution of (1) subject to $\mathbf{x} = \mathbf{S}_x \mathbf{x}_1 + \mathbf{S}_u \mathbf{u}$ is analytic, given by

$$\hat{\mathbf{u}} = (\mathbf{S}_u^\top \mathbf{Q} \mathbf{S}_u + \mathbf{R})^{-1} \mathbf{S}_u^\top \mathbf{Q} (\boldsymbol{\mu} - \mathbf{S}_x \mathbf{x}_1).$$

The residuals of this least squares solution provides information about the uncertainty of this estimate, in the form of a full covariance matrix (at control trajectory level)

$$\hat{\Sigma}^u = (\mathbf{S}_u^\top \mathbf{Q} \mathbf{S}_u + \mathbf{R})^{-1}.$$

4 iLQR optimization

Optimal control problems are defined by a cost function $\sum_{t=1}^T c(\mathbf{x}_t, \mathbf{u}_t)$ to minimize and a dynamical system $\mathbf{x}_{t+1} = \mathbf{g}(\mathbf{x}_t, \mathbf{u}_t)$ describing the evolution of a state \mathbf{x}_t driven by control commands \mathbf{u}_t during a time window of length T .

Iterative LQR (iLQR) solves such constrained nonlinear models by carrying out Taylor expansions on the cost and on the dynamical system so that a solution can be found iteratively by solving a LQR problem at each iteration. A solution in batch form can be computed by minimizing over $\mathbf{u} = [\mathbf{u}_1^\top, \mathbf{u}_2^\top, \dots, \mathbf{u}_{T-1}^\top]^\top$, yielding a series of open loop control commands \mathbf{u}_t , corresponding to a Gauss-Newton iteration scheme, see Appendix A for an overview of Newton's method for minimization. A solution can alternatively be computed in a recursive form to provide a controller with feedback gains. We consider first the batch iLQR form, as we primarily focus on the planning capability of the approach.

Batch iLQR employs a first order Taylor expansion of the dynamical system $\mathbf{x}_{t+1} = \mathbf{g}(\mathbf{x}_t, \mathbf{u}_t)$ around the point $(\hat{\mathbf{x}}_t, \hat{\mathbf{u}}_t)$, namely

$$\begin{aligned} \mathbf{x}_{t+1} &\approx \mathbf{g}(\hat{\mathbf{x}}_t, \hat{\mathbf{u}}_t) + \frac{\partial \mathbf{g}}{\partial \mathbf{x}_t} (\mathbf{x}_t - \hat{\mathbf{x}}_t) + \frac{\partial \mathbf{g}}{\partial \mathbf{u}_t} (\mathbf{u}_t - \hat{\mathbf{u}}_t) \\ \iff \Delta \mathbf{x}_{t+1} &\approx \mathbf{A}_t \Delta \mathbf{x}_t + \mathbf{B}_t \Delta \mathbf{u}_t, \end{aligned} \quad (2)$$

with error terms $\{\Delta \mathbf{x}_t = \mathbf{x}_t - \hat{\mathbf{x}}_t, \Delta \mathbf{u}_t = \mathbf{u}_t - \hat{\mathbf{u}}_t\}$, and Jacobian matrices $\{\mathbf{A}_t = \frac{\partial \mathbf{g}}{\partial \mathbf{x}_t}, \mathbf{B}_t = \frac{\partial \mathbf{g}}{\partial \mathbf{u}_t}\}$.

The cost function $c(\mathbf{x}_t, \mathbf{u}_t)$ for time step t can similarly be approximated by a second order Taylor expansion around the point $(\hat{\mathbf{x}}_t, \hat{\mathbf{u}}_t)$, namely

$$\begin{aligned} c(\mathbf{x}_t, \mathbf{u}_t) &\approx c(\hat{\mathbf{x}}_t, \hat{\mathbf{u}}_t) + \Delta \mathbf{x}_t^\top \frac{\partial c}{\partial \mathbf{x}_t} + \Delta \mathbf{u}_t^\top \frac{\partial c}{\partial \mathbf{u}_t} + \\ &\quad \frac{1}{2} \Delta \mathbf{x}_t^\top \frac{\partial^2 c}{\partial \mathbf{x}_t^2} \Delta \mathbf{x}_t + \Delta \mathbf{x}_t^\top \frac{\partial^2 c}{\partial \mathbf{x}_t \partial \mathbf{u}_t} \Delta \mathbf{u}_t + \frac{1}{2} \Delta \mathbf{u}_t^\top \frac{\partial^2 c}{\partial \mathbf{u}_t^2} \Delta \mathbf{u}_t, \end{aligned} \quad (3)$$

with gradients $\{\frac{\partial c}{\partial \mathbf{x}_t}, \frac{\partial c}{\partial \mathbf{u}_t}\}$, and Hessian matrices $\{\frac{\partial^2 c}{\partial \mathbf{x}_t^2}, \frac{\partial^2 c}{\partial \mathbf{x}_t \partial \mathbf{u}_t}, \frac{\partial^2 c}{\partial \mathbf{u}_t^2}\}$.

At a trajectory level, we denote $\mathbf{x} = [\mathbf{x}_1^\top, \mathbf{x}_2^\top, \dots, \mathbf{x}_T^\top]^\top$ the evolution of the state and $\mathbf{u} = [\mathbf{u}_1^\top, \mathbf{u}_2^\top, \dots, \mathbf{u}_{T-1}^\top]^\top$ the evolution of the control commands. The evolution of the state in (2) becomes $\Delta \mathbf{x} = \mathbf{S}_u \Delta \mathbf{u}$, see Appendix B for details.¹

The minimization problem can then be rewritten in batch form as

$$\begin{aligned} \min_{\Delta \mathbf{u}} \Delta c(\Delta \mathbf{x}, \Delta \mathbf{u}), \quad \text{s.t.} \quad \Delta \mathbf{x} &= \mathbf{S}_u \Delta \mathbf{u}, \quad \text{where} \\ \Delta c(\Delta \mathbf{x}, \Delta \mathbf{u}) &= \Delta \mathbf{x}^\top \mathbf{g}_x + \Delta \mathbf{u}^\top \mathbf{g}_u + \\ &\quad \frac{1}{2} \Delta \mathbf{x}^\top \mathbf{H}_x \Delta \mathbf{x} + \Delta \mathbf{x}^\top \mathbf{H}_{xu} \Delta \mathbf{u} + \frac{1}{2} \Delta \mathbf{u}^\top \mathbf{H}_u \Delta \mathbf{u}, \end{aligned} \quad (4)$$

with gradients $\mathbf{g}_x = \frac{\partial c}{\partial \mathbf{x}}$, $\mathbf{g}_u = \frac{\partial c}{\partial \mathbf{u}}$, and Hessian matrices $\mathbf{H}_x = \frac{\partial^2 c}{\partial \mathbf{x}^2}$, $\mathbf{H}_{xu} = \frac{\partial^2 c}{\partial \mathbf{x} \partial \mathbf{u}}$ and $\mathbf{H}_u = \frac{\partial^2 c}{\partial \mathbf{u}^2}$.

By inserting the constraint into the cost, we obtain the optimization problem

$$\begin{aligned} \min_{\Delta \mathbf{u}} \quad &\Delta \mathbf{u}^\top \mathbf{S}_u^\top \mathbf{g}_x + \Delta \mathbf{u}^\top \mathbf{g}_u + \frac{1}{2} \Delta \mathbf{u}^\top \mathbf{S}_u^\top \mathbf{H}_x \mathbf{S}_u \Delta \mathbf{u} + \\ &\Delta \mathbf{u}^\top \mathbf{S}_u^\top \mathbf{H}_{xu} \Delta \mathbf{u} + \frac{1}{2} \Delta \mathbf{u}^\top \mathbf{H}_u \Delta \mathbf{u}, \end{aligned} \quad (5)$$

which can be solved analytically by differentiating with respect to $\Delta \mathbf{u}$ and equating to zero, namely,

$$\mathbf{S}_u^\top \mathbf{g}_x + \mathbf{g}_u + \mathbf{S}_u^\top \mathbf{H}_x \mathbf{S}_u \Delta \mathbf{u} + 2 \mathbf{S}_u^\top \mathbf{H}_{xu} \Delta \mathbf{u} + \mathbf{H}_u \Delta \mathbf{u} = 0, \quad (6)$$

providing the least squares solution

$$\Delta \hat{\mathbf{u}} = (\mathbf{S}_u^\top \mathbf{H}_x \mathbf{S}_u + 2 \mathbf{S}_u^\top \mathbf{H}_{xu} + \mathbf{H}_u)^{-1} (-\mathbf{S}_u^\top \mathbf{g}_x - \mathbf{g}_u), \quad (7)$$

which can be used to update the control commands estimate at each iteration step of the iLQR algorithm.

4.1 iLQR with quadratic cost on $\mathbf{f}(\mathbf{x}_t)$

We consider a cost defined by

$$c(\mathbf{x}_t, \mathbf{u}_t) = \mathbf{f}(\mathbf{x}_t)^\top \mathbf{Q}_t \mathbf{f}(\mathbf{x}_t) + \mathbf{u}_t^\top \mathbf{R}_t \mathbf{u}_t, \quad (8)$$

where \mathbf{Q}_t and \mathbf{R}_t are weight matrices trading off task and control costs. Such cost is quadratic on $\mathbf{f}(\mathbf{x}_t)$ but non-quadratic on \mathbf{x}_t .

The cost in (3) then becomes

$$\begin{aligned} \Delta c(\Delta \mathbf{x}_t, \Delta \mathbf{u}_t) &\approx 2 \Delta \mathbf{x}_t^\top \mathbf{J}(\mathbf{x}_t)^\top \mathbf{Q}_t \mathbf{f}(\mathbf{x}_t) + 2 \Delta \mathbf{u}_t^\top \mathbf{R}_t \mathbf{u}_t + \\ &\quad \Delta \mathbf{x}_t^\top \mathbf{J}(\mathbf{x}_t)^\top \mathbf{Q}_t \mathbf{J}(\mathbf{x}_t) \Delta \mathbf{x}_t + \Delta \mathbf{u}_t^\top \mathbf{R}_t \Delta \mathbf{u}_t, \end{aligned} \quad (9)$$

¹Note that $\mathbf{S}_x \Delta \mathbf{x}_1 = \mathbf{0}$ because $\Delta \mathbf{x}_1 = \mathbf{0}$ (as we want our motion to start from \mathbf{x}_1).




Figure 1: Typical transformations involved in a manipulation task involving a robot, a vision system, a visual marker on the object, and a desired grasping location on the object.

which used

$$\frac{\partial c}{\partial \mathbf{x}_t} = 2\mathbf{J}(\mathbf{x}_t)^\top \mathbf{Q}_t \mathbf{f}(\mathbf{x}_t), \quad \frac{\partial c}{\partial \mathbf{u}_t} = 2\mathbf{R}_t \mathbf{u}_t, \quad (10)$$

$$\frac{\partial^2 c}{\partial \mathbf{x}_t^2} \approx 2\mathbf{J}(\mathbf{x}_t)^\top \mathbf{Q}_t \mathbf{J}(\mathbf{x}_t), \quad \frac{\partial^2 c}{\partial \mathbf{x}_t \partial \mathbf{u}_t} = \mathbf{0}, \quad \frac{\partial^2 c}{\partial \mathbf{u}_t^2} = 2\mathbf{R}_t. \quad (11)$$

with $\mathbf{J}(\mathbf{x}_t) = \frac{\partial \mathbf{f}(\mathbf{x}_t)}{\partial \mathbf{x}_t}$ a Jacobian matrix.

At a trajectory level, the evolution of the tracking and control weights is represented by $\mathbf{Q} = \text{blockdiag}(\mathbf{Q}_1, \mathbf{Q}_2, \dots, \mathbf{Q}_T)$ and $\mathbf{R} = \text{blockdiag}(\mathbf{R}_1, \mathbf{R}_2, \dots, \mathbf{R}_{T-1})$, respectively.

With a slight abuse of notation, we define $\mathbf{f}(\mathbf{x})$ as a vector concatenating the vectors $\mathbf{f}(\mathbf{x}_t)$, and $\mathbf{J}(\mathbf{x})$ as a block-diagonal concatenation of the Jacobian matrices $\mathbf{J}(\mathbf{x}_t)$. The minimization problem (5) then becomes

$$\min_{\Delta \mathbf{u}} 2\Delta \mathbf{x}^\top \mathbf{J}(\mathbf{x})^\top \mathbf{Q} \mathbf{f}(\mathbf{x}) + 2\Delta \mathbf{u}^\top \mathbf{R} \mathbf{u} + \Delta \mathbf{x}^\top \mathbf{J}(\mathbf{x})^\top \mathbf{Q} \mathbf{J}(\mathbf{x}) \Delta \mathbf{x} + \Delta \mathbf{u}^\top \mathbf{R} \Delta \mathbf{u}, \quad \text{s.t.} \quad \Delta \mathbf{x} = \mathbf{S}_u \Delta \mathbf{u}, \quad (12)$$

whose least squares solution is given by

$$\Delta \hat{\mathbf{u}} = (\mathbf{S}_u^\top \mathbf{J}(\mathbf{x})^\top \mathbf{Q} \mathbf{J}(\mathbf{x}) \mathbf{S}_u + \mathbf{R})^{-1} (-\mathbf{S}_u^\top \mathbf{J}(\mathbf{x})^\top \mathbf{Q} \mathbf{f}(\mathbf{x}) - \mathbf{R} \mathbf{u}), \quad (13)$$

which can be used to update the control commands estimates at each iteration step of the iLQR algorithm.

In the next sections, we show examples of functions $\mathbf{f}(\mathbf{x})$ that can rely on this formulation.

iLQR for robot manipulator

We define a manipulation task involving a set of transformations as in Fig. 1. By relying on these transformation operators, we will next describe all variables in the robot frame of reference (defined by $\mathbf{0}$, \mathbf{e}_1 and \mathbf{e}_2 in the figure).

For a manipulator controlled by joint angle velocity commands $\mathbf{u} = \dot{\mathbf{x}}$, the evolution of the system is described by $\mathbf{x}_{t+1} = \mathbf{x}_t + \mathbf{u}_t \Delta t$, with the Taylor expansion (2) simplifying to $\mathbf{A}_t = \frac{\partial \mathbf{g}}{\partial \mathbf{x}_t} = \mathbf{I}$ and $\mathbf{B}_t = \frac{\partial \mathbf{g}}{\partial \mathbf{u}_t} = \mathbf{I} \Delta t$. Similarly, a double integrator can alternatively be considered, with acceleration commands $\mathbf{u} = \ddot{\mathbf{x}}$ and states composed of both positions and velocities.




Figure 2: ReLU-like functions used in optimization costs. The derivatives of these functions are simple, providing Jacobians whose entries are either 0 or ± 1 .

For a robot manipulator, $\mathbf{f}(\mathbf{x}_t)$ in (13) typically represents the error between a reference $\boldsymbol{\mu}_t$ and the end-effector position computed by the forward kinematics function $\mathbf{f}^{ee}(\mathbf{x}_t)$. We then have

$$\mathbf{f}(\mathbf{x}_t) = \mathbf{f}^{ee}(\mathbf{x}_t) - \boldsymbol{\mu}_t, \\ \mathbf{J}(\mathbf{x}_t) = \mathbf{J}^{ee}(\mathbf{x}_t).$$

For the orientation part of the data (if considered), the Euclidean distance vector $\mathbf{f}^{ee}(\mathbf{x}_t) - \boldsymbol{\mu}_t$ is replaced by a geodesic distance measure computed with the logarithmic map $\log_{\boldsymbol{\mu}_t}(\mathbf{f}^{ee}(\mathbf{x}_t))$.

The approach can similarly be extended to target objects/landmarks with positions $\boldsymbol{\mu}_t$ and orientation matrices \mathbf{U}_t , whose columns are basis vectors forming a coordinate system, see Fig. 3. We can then define an error between the robot end-effector and an object/landmark expressed in the object/landmark coordinate system as

$$\mathbf{f}(\mathbf{x}_t) = \mathbf{U}_t^\top (\mathbf{f}^{ee}(\mathbf{x}_t) - \boldsymbol{\mu}_t), \quad (14)$$

$$\mathbf{J}(\mathbf{x}_t) = \mathbf{U}_t^\top \mathbf{J}^{ee}(\mathbf{x}_t). \quad (15)$$

Bounded joint space

The iLQR solution in (13) can be used to keep the state within a boundary (e.g., joint angle limits). We denote $\mathbf{f}(\mathbf{x}) = \mathbf{f}^{\text{cut}}(\mathbf{x})$ as the vertical concatenation of $\mathbf{f}^{\text{cut}}(\mathbf{x}_t)$ and $\mathbf{J}(\mathbf{x}) = \mathbf{J}^{\text{cut}}(\mathbf{x})$ as a diagonal concatenation of diagonal Jacobian matrices $\mathbf{J}^{\text{cut}}(\mathbf{x}_t)$. Each element i of $\mathbf{f}^{\text{cut}}(\mathbf{x}_t)$ and $\mathbf{J}^{\text{cut}}(\mathbf{x}_t)$ is defined as

$$\mathbf{f}_i^{\text{cut}}(x_{t,i}) = \begin{cases} x_{t,i} - x_i^{\min}, & \text{if } x_{t,i} < x_i^{\min} \\ x_{t,i} - x_i^{\max}, & \text{if } x_{t,i} > x_i^{\max} \\ 0, & \text{otherwise} \end{cases} \\ \mathbf{J}_{i,i}^{\text{cut}}(x_{t,i}) = \begin{cases} 1, & \text{if } x_{t,i} < x_i^{\min} \\ -1, & \text{if } x_{t,i} > x_i^{\max} \\ 0, & \text{otherwise} \end{cases}$$

where $\mathbf{f}_i^{\text{cut}}(x_{t,i})$ describes continuous ReLU-like functions for each dimension. $\mathbf{f}_i^{\text{cut}}(x_{t,i})$ is 0 inside the bounded domain and takes the signed distance value outside the boundary, see Fig. 2-left.

We can see with (10) that for $\mathbf{Q} = \frac{1}{2}\mathbf{I}$, if \mathbf{x} is outside the domain during some time steps t , $\mathbf{g}_\mathbf{x} = \frac{\partial c}{\partial \mathbf{x}} = 2\mathbf{J}^\top \mathbf{Q} \mathbf{f}$ generates a vector bringing it back to the boundary of the domain.

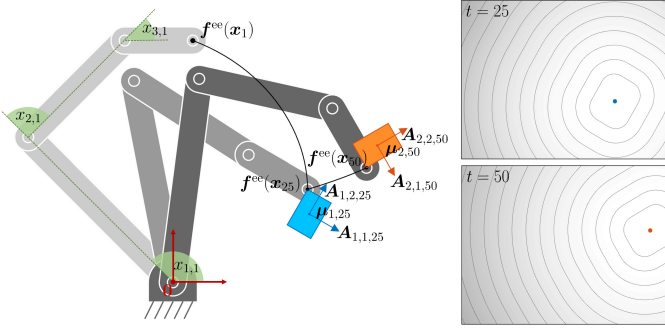


Figure 3: Example of a viapoints task in which a planar robot with 3 joints needs to sequentially reach 2 objects, with object boundaries defining the allowed reaching points on the objects surfaces. *Left*: Reaching task with two viapoints at $t = 25$ and $t = 50$. *Right*: Corresponding values of the cost function for the end-effector space at $t = 25$ and $t = 50$.

Bounded task space

The iLQR solution in (13) can be used to keep the end-effector within a boundary (e.g., end-effector position limits). Based on the above definitions, $\mathbf{f}(\mathbf{x})$ and $\mathbf{J}(\mathbf{x})$ are in this case defined as

$$\begin{aligned}\mathbf{f}(\mathbf{x}_t) &= \mathbf{f}^{\text{cut}}(\mathbf{e}(\mathbf{x}_t)), \\ \mathbf{J}(\mathbf{x}_t) &= \mathbf{J}^{\text{cut}}(\mathbf{e}(\mathbf{x}_t)) \mathbf{J}^{\text{ee}}(\mathbf{x}), \\ \text{with } \mathbf{e}(\mathbf{x}_t) &= \mathbf{f}^{\text{ee}}(\mathbf{x}).\end{aligned}$$

Reaching task with robot manipulator and prismatic object boundaries

The definition of $\mathbf{f}(\mathbf{x}_t)$ and $\mathbf{J}(\mathbf{x}_t)$ in (14) and (15) can also be extended to objects/landmarks with boundaries by defining

$$\begin{aligned}\mathbf{f}(\mathbf{x}_t) &= \mathbf{f}^{\text{cut}}(\mathbf{e}(\mathbf{x}_t)), \\ \mathbf{J}(\mathbf{x}_t) &= \mathbf{J}^{\text{cut}}(\mathbf{e}(\mathbf{x}_t)) \mathbf{U}_t^\top \mathbf{J}^{\text{ee}}(\mathbf{x}_t), \\ \text{with } \mathbf{e}(\mathbf{x}_t) &= \mathbf{U}_t^\top (\mathbf{f}^{\text{ee}}(\mathbf{x}_t) - \boldsymbol{\mu}_t),\end{aligned}$$

see also Fig. 3.

Appendices

A Newton's method for minimization

Newton's method attempts to solve $\min_x f(x)$ or $\max_x f(x)$ from an initial guess x_1 by using a sequence of second-order Taylor approximations of f around the iterates, see Fig. 4. The second-order Taylor expansion of f around x_k is

$$f(x_k + t) \approx f(x_k) + f'(x_k)t + \frac{1}{2}f''(x_k)t^2. \quad (16)$$

The next iterate $x_{k+1} = x_k + t$ is defined so as to minimize this quadratic approximation in t . If the second




Figure 4: Newton's method for minimization, starting from an initial estimate x_1 and converging to a local minimum (red point) after 5 iterations.

derivative is positive, the quadratic approximation is a convex function of t , and its minimum can be found by setting the derivative to zero. Since

$$\frac{d}{dt} \left(f(x_k) + f'(x_k)t + \frac{1}{2}f''(x_k)t^2 \right) = f'(x_k) + f''(x_k)t, \quad (17)$$

the minimum is achieved for

$$t = -\frac{f'(x_k)}{f''(x_k)}. \quad (18)$$

Newton's method thus performs the iteration

$$x_{k+1} = x_k - \frac{f'(x_k)}{f''(x_k)}. \quad (19)$$

The geometric interpretation of Newton's method is that at each iteration, it amounts to the fitting of a paraboloid to the surface of $f(x)$ at x_k , having the same slopes and curvature as the surface at that point, and then proceeding to the maximum or minimum of that paraboloid. Note that if f happens to be a quadratic function, then the exact extremum is found in one step. Note also that Newton's method is often modified to include a step size (e.g., estimated with line search).

The multidimensional case similarly provides

$$\mathbf{x}_{k+1} = \mathbf{x}_k - \mathbf{H}(\mathbf{x}_k)^{-1} \mathbf{g}(\mathbf{x}_k), \quad (20)$$

with \mathbf{g} and \mathbf{H} the gradient and Hessian matrix of f (vector and square matrix, respectively).

Gauss-Newton algorithm

The Gauss-Newton algorithm is a special case of Newton's method in which the cost is quadratic (sum of squared function values), with $f(\mathbf{x}) = \sum_{i=1}^R r_i^2(\mathbf{x}) = \mathbf{r}^\top \mathbf{r}$, and by ignoring the second-order derivative terms of the Hessian. The gradient and Hessian can in this case be computed with

$$\mathbf{g} = 2\mathbf{J}_r^\top \mathbf{r}, \quad \mathbf{H} \approx 2\mathbf{J}_r^\top \mathbf{J}_r, \quad (21)$$

where $\mathbf{J}_r \in \mathbb{R}^{R \times D}$ is the Jacobian matrix of $\mathbf{r} \in \mathbb{R}^R$. The update rule then becomes

$$\mathbf{x}_{k+1} = \mathbf{x}_k - (\mathbf{J}_r^\top(\mathbf{x}_k) \mathbf{J}_r(\mathbf{x}_k))^{-1} \mathbf{J}_r^\top(\mathbf{x}_k) \mathbf{r}(\mathbf{x}_k) \quad (22)$$

$$= \mathbf{x}_k - \mathbf{J}_r^\dagger(\mathbf{x}_k) \mathbf{r}(\mathbf{x}_k). \quad (23)$$

B System dynamics at trajectory level

The evolution of a system $\mathbf{x}_{t+1} = g(\mathbf{x}_t, \mathbf{u}_t)$ (expressed in discrete form), can be approximated by the linear system

$$\mathbf{x}_{t+1} = \mathbf{A}_t(\mathbf{x}_t, \mathbf{u}_t) \mathbf{x}_t + \mathbf{B}_t(\mathbf{x}_t, \mathbf{u}_t) \mathbf{u}_t, \quad \forall t \in \{1, \dots, T\}$$

with states $\mathbf{x}_t \in \mathbb{R}^D$ and control commands $\mathbf{u}_t \in \mathbb{R}^d$.

With the above linearization, we can express all states \mathbf{x}_t as an explicit function of the initial state \mathbf{x}_1 . By writing

$$\begin{aligned} \mathbf{x}_2 &= \mathbf{A}_1 \mathbf{x}_1 + \mathbf{B}_1 \mathbf{u}_1, \\ \mathbf{x}_3 &= \mathbf{A}_2 \mathbf{x}_2 + \mathbf{B}_2 \mathbf{u}_2 = \mathbf{A}_2(\mathbf{A}_1 \mathbf{x}_1 + \mathbf{B}_1 \mathbf{u}_1) + \mathbf{B}_2 \mathbf{u}_2, \\ &\vdots \\ \mathbf{x}_T &= \left(\prod_{t=1}^{T-1} \mathbf{A}_{T-t} \right) \mathbf{x}_1 + \left(\prod_{t=1}^{T-2} \mathbf{A}_{T-t} \right) \mathbf{B}_1 \mathbf{u}_1 + \\ &\quad \left(\prod_{t=1}^{T-3} \mathbf{A}_{T-t} \right) \mathbf{B}_2 \mathbf{u}_2 + \dots + \mathbf{B}_{T-1} \mathbf{u}_{T-1}, \end{aligned}$$

in a matrix form, we get an expression of the form $\mathbf{x} = \mathbf{S}_x \mathbf{x}_1 + \mathbf{S}_u \mathbf{u}$, with

$$\underbrace{\begin{bmatrix} \mathbf{x}_1 \\ \mathbf{x}_2 \\ \mathbf{x}_3 \\ \vdots \\ \mathbf{x}_T \end{bmatrix}}_{\mathbf{x}} = \underbrace{\begin{bmatrix} \mathbf{I} \\ \mathbf{A}_1 \\ \mathbf{A}_2 \mathbf{A}_1 \\ \vdots \\ \prod_{t=1}^{T-1} \mathbf{A}_{T-t} \end{bmatrix}}_{\mathbf{S}_x} \mathbf{x}_1 + \underbrace{\begin{bmatrix} \mathbf{0} & \mathbf{0} & \dots & \mathbf{0} \\ \mathbf{B}_1 & \mathbf{0} & \dots & \mathbf{0} \\ \mathbf{A}_2 \mathbf{B}_1 & \mathbf{B}_2 & \dots & \mathbf{0} \\ \vdots & \vdots & \ddots & \vdots \\ \left(\prod_{t=1}^{T-2} \mathbf{A}_{T-t} \right) \mathbf{B}_1 & \left(\prod_{t=1}^{T-3} \mathbf{A}_{T-t} \right) \mathbf{B}_2 & \dots & \mathbf{B}_{T-1} \end{bmatrix}}_{\mathbf{S}_u} \underbrace{\begin{bmatrix} \mathbf{u}_1 \\ \mathbf{u}_2 \\ \vdots \\ \mathbf{u}_{T-1} \end{bmatrix}}_{\mathbf{u}},$$

where $\mathbf{S}_x \in \mathbb{R}^{dT \times d}$, $\mathbf{x}_1 \in \mathbb{R}^d$, $\mathbf{S}_u \in \mathbb{R}^{dT \times d(T-1)}$ and $\mathbf{u} \in \mathbb{R}^{d(T-1)}$.

Transfer matrices for single integrator

A single integrator is simply defined as $\mathbf{x}_{t+1} = \mathbf{x}_t + \mathbf{u}_t \Delta t$, corresponding to $\mathbf{A}_t = \mathbf{I}$ and $\mathbf{B}_t = \mathbf{I} \Delta t$, $\forall t \in \{1, \dots, T\}$, and transfer matrices $\mathbf{S}_x = \mathbf{1}_T \otimes \mathbf{I}_D$, and $\mathbf{S}_u = \begin{bmatrix} \mathbf{0}_{D, D(T-1)} \\ \mathbf{L}_{T-1, T-1} \otimes \mathbf{I}_D \Delta t \end{bmatrix}$, where \mathbf{L} is a lower triangular matrix and \otimes is the Kronecker product operator.